

Signals and Signal Handling - Part 1

What Is Signals?

- A primitive way of doing IPC and most widely known UNIX facility.
 - Be used to inform processes of *asynchronous* events.
 - Posted by one process and received by another or the same process.
 - An asynchronous event either terminates a process or is simply being ignored.
 - Signal handling
 - Default (SIG_DFL)
 - Ignored (SIG_IGN)
 - User-defined

Signal Generation

- A signal is generated when (not a complete list):
 - A hardware exception occurs.
 - Interrupt or quit from control terminal.
 - An alarm timer expires.
 - A call to kill().
 - Termination of a child process.

Signal Usage

- Signals can be used:
 - Intraprocess
 - With the same user ID
 - Interprocesses
 - Between kernel to any process.

Signal States

- A signal is said to be:
 - *Generated* when the event that causes the signal occurs.
 - *Delivered* when the action for a signal is taken.
 - *Pending* during the time between the generation of the signal and its delivery.
 - *Blocked* if unable to deliver due to a signal mask bit being set for the signal.

Signal Disposition

- Default action (SIG_DFL)
 - Termination in general.
- Ignored (SIG_IGN)
 - Never posted to the process.
- User-defined action
 - Needs a user-defined *signal handler*, or *signal-catching function*.
 - Most signals can be caught, or ignored except SIGKILL and SIGSTOP.

Linux Signals (1)

SIGHUP	1 /* Hangup (POSIX). terminate w/ core */
SIGINT	2 /* Interrupt (ANSI). terminate */
SIGQUIT	3 /* Quit (POSIX). terminate w/ core */
SIGILL	4 /* Illegal instruction (ANSI). terminate w/ core */
SIGTRAP	5 /* Trace trap (POSIX). terminate w/ core */
SIGABRT	6 /* Abort (ANSI). terminate w/ core */
SIGIOT	6 /* IOT trap (4.2 BSD). terminate w/ core */
SIGBUS	7 /* BUS error (4.2 BSD). terminate w/ core */
SIGFPE	8 /* Floating-point exception (ANSI). terminate w/ core */
SIGKILL	9 /* Kill, unmaskable (POSIX). terminate */
SIGUSR1	10 /* User-defined signal 1 (POSIX). terminate */

Linux Signals (2)

SIGSEGV	11 /* Segmentation violation(ANSI). terminate w/ core */
SIGUSR2	12 /* User-defined signal 2 (POSIX). terminate */
SIGPIPE	13 /* Broken pipe (POSIX). terminate */
SIGALRM	14 /* Alarm clock (POSIX). terminate */
SIGTERM	15 /* Termination (ANSI). terminate */
SIGSTKFLT	16 /* Stack fault. terminate w/ core */
SIGCLD	SIGCHLD /* Same as SIGCHLD (System V). */
SIGCHLD	17 /* Child status has changed(POSIX). ignore */
SIGCONT	18 /* Continue (POSIX). continue/ignore */
SIGSTOP	19 /* Stop, unmaskable (POSIX). stop process */
SIGTSTP	20 /* Keyboard stop (POSIX). stop process */
SIGTTIN	21 /* Background read from tty (POSIX). stop process */
SIGTTOU	22 /* Background write to tty (POSIX). stop process */

Linux Signals (3)

SIGURG	23	/* Urgent condition on socket (BSD). ignore */
SIGXCPU	24	/* CPU limit exceeded (BSD). terminate w/ core */
SIGXFSZ	25	/* File size limit exceeded (BSD). terminate w/ core */
SIGVTALRM	26	/* Virtual alarm clock (BSD). terminate */
SIGPROF	27	/* Profiling alarm clock (BSD). terminate */
SIGWINCH	28	/* Window size change (BSD,Sun). ignore */
SIGPOLL	SIGIO	/* Pollable event occurred (System V). terminate */
SIGIO	29	/* I/O now possible (BSD). terminate/ignore */
SIGPWR	30	/* Power failure restart (System V). ignore */
SIGUNUSED	31	

Signal Sets

- Signal sets are one of the main parameters passed to system calls that deal with signals
- A list of signals you want to do something with
- To manipulate signal sets, a new data type known as `sigset_t` with the following five predefined functions is specified in POSIX.1:
 - `sigemptyset()`
 - `sigfillset()`
 - `sigaddset()`
 - `sigdelset()`
 - `sigismember()`

Signal Set (1)

- `#include <signal.h>`
`int sigemptyset(sigset_t *set);`
 - Initializes the signal set given by `set` to empty, with all signals excluded from the set.
 - Return 0 on success and -1 on error.
- `int sigfillset(sigset_t *set);`
 - Initializes `set` to full, including all signals.
 - Return 0 on success and -1 on error.

Signal Set (2)

- `int sigaddset(sigset_t *set, int signum);`
 - Add signal *signum* from *set*.
 - Return 0 on success and -1 on error.
- `int sigdelset(sigset_t *set, int signum);`
 - Delete signal *signum* from *set*.
 - Return 0 on success and -1 on error.
- `int sigismember(const sigset_t *set, int signum);`
 - Tests whether *signum* is a member of *set*.
 - Returns 1 if *signum* is a member of *set*, 0 if *signum* is not a member, and -1 on error.

Example #1: Signal Sets

```
#include <stdio.h>
#include <signal.h>

sigset_t mask1, mask2;
..
..
/* create empty set */
sigemptyset(&mask1);

/* add signal */
sigaddset(&mask1, SIGINT);
sigaddset(&mask1, SIGQUIT);
..

/* create full set */
sigfillset(&mask2);

/* remove signal */
sigdelset(&mask2, SIGCHLD);
..
..
```

sigaction (1)

- `#include <signal.h>`

`int sigaction(int signum, const struct sigaction *act, struct sigaction *oldact);`

- Change the action taken by a process on receipt of a specific signal.
- Return 0 on success and -1 on error.
- *signum*
 - Specifies the signal and can be any valid signal except SIGKILL and SIGSTOP.
- If *act* is non-null, the new action for signal *signum* is installed from *act*.
- If *oldact* is non-null, the previous action is saved in *oldact*.

sigaction (2)

– sigaction structure

```
struct sigaction {  
    void (*sa_handler)(int);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_sigaction)(int, siginfo_t *, void);  
};
```

– sa_handler

- Specifies the action to be associated with *signum* and may be SIG_DFL for the default action.
- SIG_IGN to ignore this signal, or a pointer to a signal handling function.

sigaction (3)

- sa_mask
 - Gives a mask of signals which should be blocked during execution of the signal handler.
- sa_flags
 - Specifies a set of flags which modify the behavior of the signal handling process.
 - It is formed by the bitwise OR of zero or more of the following:
 - SA_NOCLDSTOP
 - » If *signum* is SIGCHLD, do not receive notification when child processes stop.
 - SA_RESETHAND
 - » Restore the signal action to the default state once the signal handler has been called.
- sa_sigaction(int, siginfo_t *, void *)
 - If sa_flags is set to SA_SIGINFO, extra information will be passed to the signal handler. In this case, sa_sigaction() is used

Example #2: catching SIGINT (1)

```
#include <stdio.h>
#include <signal.h>

void catchint (int signo)
{
    printf ("□nCATCHINT: signo=%d□n", signo);
    printf("CATCHINT: returning□n□n");
}

main()
{
    static struct sigaction act;

    act.sa_handler = catchint;

    sigfillset(&(act.sa_mask));

    sigaction(SIGINT, &act, NULL);
```

Example #2: catching SIGINT (2)

```
printf ("sleep call #1\n");  
sleep (1);  
printf ("sleep call #2\n");  
sleep (1);  
printf ("sleep call #3\n");  
sleep (1);  
printf ("sleep call #4\n");  
sleep (1);  
printf ("Exiting\n");  
exit (0);  
}
```

% a.out
sleep call #1
sleep call #2

CATCHINT: signo=2
CATCHINT: returning

sleep call #3
sleep call #4
Exiting

Example #3: ignoring SIGINT

- Just replace the following line in the example #2 program

- `act.sa_handler = catchint;`

With:

- `act.sa_handler = SIG_IGN;`

And then call `sigaction(SIGINT, &act, NULL)`

Example #4: restoring a previous action

```
#include <signal.h>
```

```
static struct sigaction act, oact;
```

```
/* save the old action for SIGTERM */  
sigaction(SIGTERM, NULL, &oact);
```

```
/* set new action for SIGTERM */  
sigaction(SIGTERM, &act, NULL);
```

```
/* do the work here.... */
```

```
/* now restore the old action */  
sigaction(SIGTERM, &oact, NULL);
```

Example #5: graceful exit

- Suppose a program uses a temporary workfile

```
/* exit from program gracefully */  
#include <stdio.h>  
#include <stdlib.h>  
void g_exit(int s)  
{  
    unlink("tempfile");  
    fprintf(stderr, "Interrupted – exiting\n");  
    exit(1)  
}
```

```
/* in somewhere */  
static struct sigaction act;  
act.sa_handler = g_exit;  
sigaction(SIGTINT, &act, NULL);
```

Signal Handler

- `void (*signal(int signo, void (*handler)(int)))(int)`
 - Signal handler can be set by user process.
- `signal()` is said to be unreliable.
 - Signals can get lost.
- Further superseded by `sigaction()` in the latest implementations of various versions of UNIX systems.

signal (1)

- `#include <signal.h>`

`void (*signal(int signum, void (*handler)(int)))(int);`

- Installs a new signal handler for the signal with number *signum*.
- The signal handler is set to *handler* which may be a user specified function, or one of the following:
 - SIG_IGN: Ignore the signal.
 - SIG_DFL: Reset the signal to its default behavior.

signal (2)

- The integer argument that is handed over to the signal handler routine is the *signal number*.
- It is possible to use one signal handler for several signals.
- Return value
 - The previous value (address) of the signal handler, or SIG_ERR on error.

Example #6: signal (1)

```
#include <stdio.h>
#include <signal.h>
#include <sys/types.h>

static void sig_usr(int signo)
{
    if (signo == SIGUSR1)
        printf("received SIGUSR1\n");
    else if (signo == SIGUSR2)
        printf("received SIGUSR2\n");
    else {
        fprintf(stderr, "received signal %d\n", signo);
        fflush(stderr);
        abort();
    }
    return;
}
```

Example #6: signal (2)

```
int main(void)
{
    if (signal(SIGUSR1, sig_usr) == SIG_ERR) {
        perror("can't catch SIGUER1");
        exit(1);
    }
    if (signal(SIGUSR2, sig_usr) == SIG_ERR) {
        perror("can't catch SIGUER2");
        exit(1);
    }

    for(;;)
        pause();
}
```

Example #6: Sample Run

```
% a.out &  
[1] 4720  
% kill -USR1 4720  
received SIGUSR1  
% kill -USR2 4720  
received SIGUSR2  
% kill 4720  
[1] + Terminated a.out &
```

Signal Mask

- Be used to block signal delivery.
 - A blocked signal depends on the recipient process to unblock and handle it accordingly.
- A signal mask may be implemented using an integer.
 - Positional -- each bit corresponds to one signal.
 - Bit 1's -- the corresponding signals are being blocked.
- A process may query or change its signal mask by a call to `sigprocmask()`.

sigprocmask (1)

- `int sigprocmask(int how, const sigset_t *set, sigset_t *oldset);`
 - Change the list of currently blocked signals.
 - Return 0 on success and -1 on error.
 - *oldset*
 - If non-null, the previous value of the signal mask is stored in *oldset*.
 - It is not possible to block SIGKILL or SIGSTOP with the sigprocmask call.
 - If *set* is NULL, *how* is ignored and the current value of the signal mask is returned by *oldset*.

sigprocmask (2)

- *how*

- SIG_BLOCK
 - The set of blocked signals is the *union* of the current set and the *set* argument.
- SIG_UNBLOCK
 - The signals in *set* are removed from the current set of blocked signals.
 - It is legal to attempt to unblock a signal which is not blocked.
- SIG_SETMASK
 - The set of blocked signals is *set* to the argument *set*.

Example #7: sigprocmask (1)

```
/* signal blocking -- sigprocmask의 사용예를 보인다. */
```

```
#include <signal.h>
```

```
main()
```

```
{
```

```
    sigset_t set1, set2;
```

```
    /* 시그널 집합을 완전히 채운다. */
```

```
    sigfillset (&set1);
```

```
    /* SIGINT와 SIGQUIT를 포함하지 않는 시그널 집합을 생성한다. */
```

```
    sigfillset (&set2);
```

```
    sigdelset (&set2, SIGINT);
```

```
    sigdelset (&set2, SIGQUIT);
```

Example #7: sigprocmask (2)

```
/* 중대하지 않은 코드를 수행 ... */

/* 봉쇄를 설정한다. */
sigprocmask(SIG_SETMASK, &set1, NULL);

/* 극도로 중대한 코드를 수행한다. */

/* 하나의 봉쇄를 제거한다. */
sigprocmask(SIG_UNBLOCK, &set2, NULL);

/* 덜 중대한 코드를 수행한다 ... */

/* 모든 시그널 봉쇄를 제거한다. */
sigprocmask(SIG_UNBLOCK, &set1, NULL);
}
```


Example #8: sigprogmask

```
#include <stdio.h>
#include <signal.h>
#include <sys/types.h>

void pr_mask(const char *str)
{
    sigset_t sigset;

    if (sigprocmask(0, NULL, &sigset)
        < 0) {
        perror("sigprocmask error");
        exit(1);
    }

    printf("s%", str);
    if (sigismember(&sigset, SIGINT))
        printf("SIGINT");
    if (sigismember(&sigset, SIGQUIT))
        printf("SIGQUIT");
    if (sigismember(&sigset, SIGUSR1))
        printf("SIGUSR1");
    if (sigismember(&sigset, SIGALRM))
        printf("SIGALRM");

    /* remaining signals can go here */

    printf("\n");
}
```

sigpending

- `#include <signal.h>`

`int sigpending(sigset_t *set);`

- Examine the pending signals.
- The signal mask of pending signals is stored in *set*.
- Return 0 on success and -1 on error.

Example #9: sigpending (1)

```
#include <stdio.h>
#include <signal.h>
#include <sys/types.h>

static void sig_quit(int);

int main(void)
{
    sigset_t newmask, oldmask, pendmask;

    if (signal(SIGQUIT, sig_quit) == SIG_ERR) {
        perror("can't catch SIGQUIT");
        exit(1);
    }

    /* block SIGQUIT and save current signal mask */
    sigemptyset(&newmask);
    sigaddset(&newmask, SIGQUIT);
```

Example #9: sigpending (2)

```
if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) {
    perror("SIG_BLOCK error");
    exit(1);
}
sleep(5);

if (sigpending(&pendmask) < 0) {
    perror("sigpending error");
    exit(1);
}

if (sigismember(&pendmask, SIGQUIT))
    printf("\nSIGQUIT pending\n");

/* reset signal mask which unblock SIGQUIT */
if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) {
    perror("SIG_SETMASK error");
    exit(1);
}
printf("SIGQUIT unblocked\n");
```

Example #9: sigpending (3)

```
/* SIGQUIT here will terminate with core file */
sleep(5);
exit(0);
}

static void sig_quit(int signo)
{
    printf("caught SIGQUIT\n");

    if (signal(SIGQUIT, SIG_DFL) == SIG_ERR) {
        perror("can't reset SIGQUIT");
        exit(1);
    }
    return;
}
```

Example #9: Sample Run

% a.out

^\ *generate signal once (before 5 seconds are up)*

SIGQUIT pending *after return from sleep*

caught SIGQUIT *in signal handler*

SIGQUIT unblocked *after return from sigprocmask*

^\Quit *generate (pending) signal again*

% a.out

^\^\^\^\^\^\^\^\^\^\^\^\^\^\^\^\ *generate signal 10 times before 5 seconds are up*

SIGQUIT pending *after return from sleep*

caught SIGQUIT *signal is generated only once*

SIGQUIT unblocked *after return from sigprocmask*

^\Quit *generate (pending) signal again*